

Chapter 7

Multiple-Layer Neural Networks

7.1 Learning Objectives

The previous chapter addressed the basic concepts of neural networks with one hidden layer. These concepts form a foundation of deep learning machines. In this chapter, we will extend these concepts to general multiple-layer neural networks, called *deep neural networks*, which consist of many hidden layers. Multiple-layer neural networks or deep learning are often used to recognize complicated patterns, such as face recognition and nature language process, such as speech recognition, language translation, and speech-text transcription. To make it concrete, we use an example, image classification, to demonstrate a deep NN development which is close to practical applications. Upon the completion of this chapter, one should be able to

- Understand and use ReLU activation function.
- To initialize parameters of neural network.
- Apply gradient descent to multiple-layer neural network training.
- Gain some programming skills, such as image reading and display, accuracy computing.
- Develop neural networks with any size from the scratch.

7.2 Representation of Multiple-Layer Neural Networks

7.2.1 Architecture

A deep neural network consists of multiple hidden layers with multiple units in each layer. For instance, a fully connected 4-hidden-layer neural network with 4 units in the output layer is shown in Fig.1. The input x has 8 features. There are 9 units in each hidden layer. Please note that different hidden layers may have different numbers of units. The network delivers 4 outputs.




Fig.1 multiple layer neural network




Fig. 2 diagram of layer l .

Now let's focus on one particular layer. Let L be the total number of layers (not including the input layer, input layer is called layer [0], just input data). Thus, layer l , $l=1,2,\dots,L-1$, is hidden layer, and layer L is the output layer. In general, layer l can be defined by weight matrix $W^{[l]}$ and bias vector $b^{[l]}$, and an activation function $g^{[l]}(\cdot)$ for each unit. The input and output of layer l are denoted by vectors $a^{[l-1]}$ and $a^{[l]}$, respectively. The block diagram representation of layer l is shown in Fig.2. The shapes of $W^{[l]}$ and $b^{[l]}$ are defined by the number of units ($n^{[l]}$ and $n^{[l-1]}$) in layer l and layer $l-1$.

Now let's zoom in layer l with parameters $W^{[l]}$ and $b^{[l]}$. The shape of $W^{[l]}$ is $(n^{[l]}, n^{[l-1]})$, and shape of $b^{[l]}$ is $(n^{[l]}, 1)$. The i th row in $W^{[l]}$ and $b^{[l]}$ is responsible for the input of the i th unit in the layer. The relationship between the input and output in layer l can be described as a forward propagation

$$z^{[l]} = W^{[l]}a^{[l-1]} + b^{[l]} \quad (1.a)$$

Check shape consistence for (1.a): $(n^{[l]}, 1) = (n^{[l]}, n^{[l-1]}) \times (n^{[l-1]}, 1) + (n^{[l-1]}, 1)$

$$a^{[l]} = g^{[l]}(z^{[l]}) \quad (1.b)$$

Check shape consistence for (1.b): $(n^{[l]}, 1) = g^{[l]}(n^{[l]}, 1)$. Please note that (1.a) and (1.b) is based on one data example.

7.2.2 Forward propagation and backward propagation

The data flow, shown in Fig.3, from the input layer to the output layer is referred as forward propagation, which involves computing of (1.a) and (1.b) for all layers. The purpose of calculating forward propagation is two-fold: 1) to deliver quantities needed for derivative computing during the training process; 2) to predict the output for an input x during the inference for new data examples. The backward propagation will provide all derivatives of cost function with respect to parameters, which are needed in the gradient descent algorithm. The gradient descent algorithm updates the parameters until the optimal parameters have been found according to a stopping criterion.

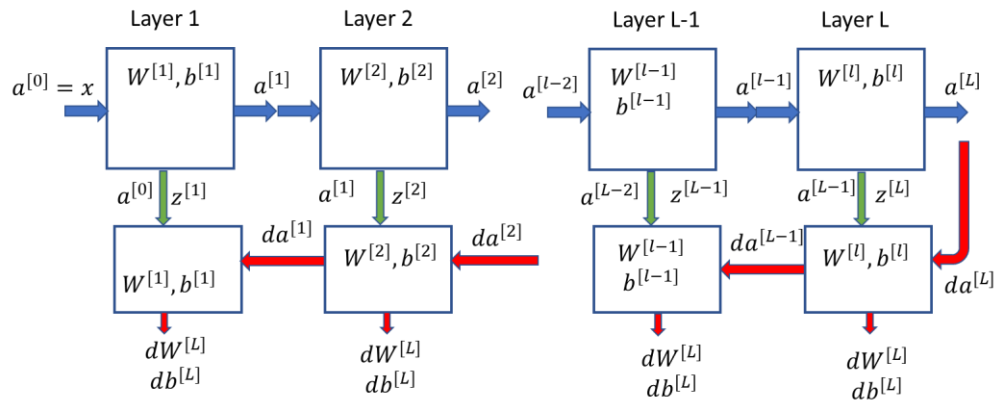


Fig.3 forward and backward block diagram

In Fig.3, the top row of boxes shows the forward propagation, i.e. the neural network, and the bottom row the backward propagation for derivative computation. In the previous chapter, we derived the equations in forward propagation and backward propagation as follows.

1) Forward propagation for layer l

Input: $a^{[l-1]}$

Output: $a^{[l]}, z^{[l]}$ (saved to cache)

For one data example x

$$a^{[0]} = x$$

$$z^{[l]} = W^{[l]}a^{[l-1]} + b^{[l]}$$

$$a^{[l]} = g^{[l]}(z^{[l]})$$

For m data examples (vectorized)

(shape)

$$A^{[0]} = X$$

(n_x, m)

$$Z^{[l]} = W^{[l]}A^{[l-1]} + b^{[l]}$$

$$(n^{[l]}, m) = (n^{[l]}, n^{[l-1]}) \times (n^{[l-1]}, m) + (n^{[l-1]}, 1)$$

$$A^{[l]} = g^{[l]}(Z^{[l]})$$

$(n^{[l]}, m)$

2) Backward propagation for layer l

Input: $da^{[l]}$ (from forward propagation), $a^{[l-1]}, z^{[l]}$ (from cache)

Output: $da^{[l-1]}, dW^{[l]}, db^{[l]}$

For one data example x

(shape)

$$da^{[L]} = -\frac{y}{a^{[L]}} + \frac{1-y}{1-a^{[L]}} \quad (\text{assume sigmoid function in the output layer})$$

$$da^{[L]} = \frac{dL(\hat{y}, y)}{da^{[L]}} = -\frac{y}{a^{[L]}} \quad (\text{assume softmax in the output layer})$$

$$dz^{[L]} = a^{[L]} - y \quad (\text{for either sigmoid or softmax})$$

$$dz^{[l]} = da^{[l]} * g^{[l]'}(z^{[l]})$$

$$(n^{[l]}, 1) = (n^{[l]}, 1) * (n^{[l]}, 1)$$

$$dW^{[l]} = dz^{[l]}a^{[l-1]T}$$

$$(n^{[l]}, n^{[l-1]}) = (n^{[l]}, 1) \times (n^{[l-1]}, 1)^T$$

$$db^{[l]} = dz^{[l]}$$

$$(n^{[l]}, 1) = (n^{[l]}, 1)$$

$$da^{[l-1]} = W^{[l]T} dz^{[l]}$$

$$(n^{[l-1]}, 1) = (n^{[l]}, n^{[l-1]})^T \times (n^{[l]}, 1)$$

For m data examples (vectorized)

$$dA^{[L]} = -\frac{Y}{A^{[L]}} + \frac{1-Y}{1-A^{[L]}}$$

$(1, m)$ for sigmoid function at layer L

$$dA^{[L]} = -\frac{Y}{A^{[L]}}$$

$(n^{[L]}, m)$ for softmax function at layer L

$$dZ^{[L]} = A^{[L]} - Y$$

$(n^{[L]}, m)$ for either sigmoid or softmax

$$dz^{[l]} = da^{[l]} * g^{[l]'}(z^{[l]})$$

$$(n^{[l]}, m) = (n^{[l]}, m) * (n^{[l]}, m)$$

$$dW^{[l]} = dz^{[l]}A^{[l-1]T} * \left(\frac{1}{m}\right)$$

$$(n^{[l]}, n^{[l-1]}) = (n^{[l]}, m) \times (n^{[l-1]}, m)^T$$

$$db^{[l]} = \left(\frac{1}{m}\right) np.sum(dz^{[l]}, axis = 1, keepdims = True)$$

$$(n^{[l]}, 1) = (n^{[l]}, 1)$$

$$da^{[l-1]} = W^{[l]T} dz^{[l]}$$

$$(n^{[l-1]}, 1) = (n^{[l]}, n^{[l-1]})^T \times (n^{[l]}, 1)$$

Attention should be paid to the differences when different activation functions (sigmoid for binary classification or softmax for multiple classification) are applied in the output layer. In the case of softmax, Y is encoded as one-hot code matrix with a shape of $(n^{[L]}, m)$. Although $dA^{[L]}$ has different equations for sigmoid and softmax functions, $dZ^{[L]} = A^{[L]} - Y$ is the same for both (see (6.35) and (6.56) for the reason).

7.3 L-layer Neural Network for Image Classification: A Tutorial

In this section, we will a neural network for image classification. In the dataset, some images include a cat while others don't. Our trained neural network will recognize whether an image has a cat in it. The image size is $(64 \times 64 \times 3)$, i.e., 64-by-64 pixels, and 3 values (for red, blue and green) at each pixel. An image is converted into a vector of 12288 elements as the input to the neural network. The hyperparameters and parameters of the neural network include:

- Learning rate: α
- The number of iterations for parameter update
- The number of layers, L .
- The number of units for each layer: $n^{[1]}, n^{[2]}, \dots, n^{[L]}$
- Choice of activation functions: ReLu for hidden layers, sigmoid for output layer
- Learned parameters: $W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}, \dots, W^{[L]}, b^{[L]}$

The architecture of the neural network is illustrated by Fig. 4, where ReLu is a widely used activation function for computer vision, defined by

$$ReLu(z) = \begin{cases} z & z \geq 0 \\ 0 & otherwise \end{cases} \quad (2)$$

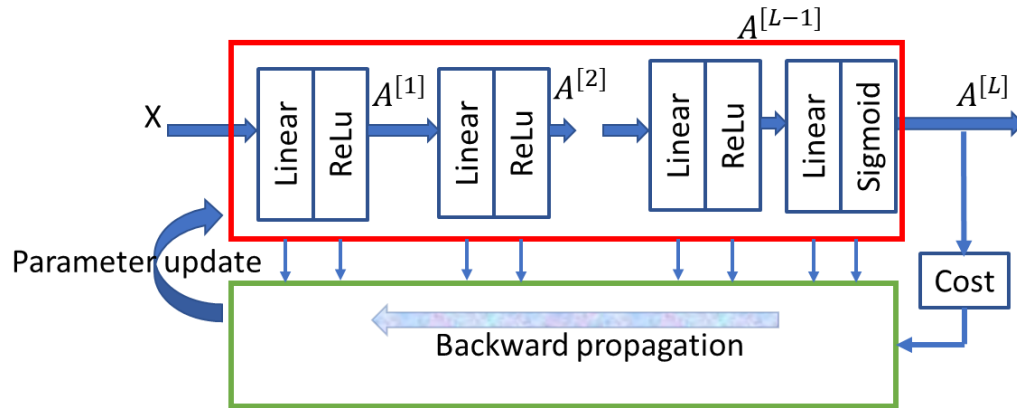


Fig.4 L-layer neural network with ReLu in hidden layers

The development process includes the following steps:

- 1) Initialize the parameters for an L-layer neural network.
- 2) Implement the forward propagation module (shown in the red box in Fig.4).

- Complete the LINEAR part of a layer's forward propagation step (resulting in $Z[l]$).
 - Complete ACTIVATION function (relu/sigmoid).
 - Combine the previous two steps into a LINEAR_ACTIVATION forward function.
 - Stack the LINEAR_RELU block L-1 time (for layers 1 through L-1) and add a LINEAR_SIGMOID block at the end (for the final layer L). This results in a new `L_model_forward` function.
- 3) Compute the loss.
 - 4) Implement the backward propagation module (denoted in the green box in Fig.4).
 - Complete the LINEAR part of a layer's backward propagation step.
 - Compute gradient of the ACTIVATE function (relu_backward/sigmoid_backward)
 - Combine the previous two steps into a LINEAR_ACTIVATION backward function.
 - Stack LINEAR_RELU backward block L-1 times and add LINEAR_SIGMOID backward in a new `L_model_backward` function
 - 5) Finally update the parameters.

7.3.1 Implementation of block functions

We develop the basic blocks/functions for the DNN model, which were mentioned above. All the functions are included in the file `dnn_utils.py`. To utilize these functions, a top-level file needs to import the package `dnn_utils` at the beginning.

- 1) Functions `sigmoid(Z)`, `relu(Z)`, `relu_backward(dA,Z)`, and `sigmoid_backward(dA,Z)`
`sigmoid(Z)`, `relu(Z)`: compute the value of the function, and return both function value and Z .
`relu_backward(dA,Z)`: compute $\frac{dJ}{dZ} = \frac{dJ}{dA} \frac{dA}{dZ}$ (denoted as: $dZ = dA \frac{dA}{dZ}$, where $\frac{dA}{dZ} = 1, Z > 0$; $\frac{dA}{dZ} = 0, Z < 0$), and return the derivative dZ .
`sigmoid_backward(dA,Z)`: compute $\frac{dJ}{dZ} = \frac{dJ}{dA} \frac{dA}{dZ}$ (denoted as: $dZ = dA \frac{dA}{dZ}$, where $\frac{dA}{dZ} = \text{sigmoid}(Z) * (1 - \text{sigmoid}(Z))$), and return the derivative dZ
- 2) Initializations of parameters
`initialize_parameters_deep(layer_dims)`: initialize the parameters for L-layer neural network. The input `layer_dims` is an array that defines the number of units in each layer. Both two functions output a dictionary that include parameters W and b
 Note: Scaling of the initial parameter is important. For the 2-layer nn, both `/np.sqrt(n_x)` and `*0.01` work. But for L-layer nn, `*0.01` does not work (demonstrated by the later image classification example). Thus, the parameters are scaled by `/ np.sqrt(layer_dims[l-1])`, not `*0.01`.

```
parameters['W' + str(l)] = np.random.randn(layer_dims[l],layer_dims[l-1])/ np.sqrt(layer_dims[l-1])
```

- 3) `linear_forward(A,W,b)`

$$Z^{[l]} = W^{[l]}A^{[l-1]} + b^{[l]}$$

`linear_forward(A, W, b)`: compute $Z = \text{np.dot}(W,A)+b$, return Z and `cache=(A,W,b)`

- 4) `linear_activation_forward(A_prev, W, b, activation)`: combine the `linear_forward(A,W,b)` and `sigmoid(Z)` or `relu(Z)` functions, and compute

$$Z^{[l]} = W^{[l]}A^{[l-1]} + b^{[l]}$$

$$A^{[l]} = g^{[l]}(Z^{[l]})$$

return $A = A^{[l]}$, and cache $(A^{[l-1]}, W^{[l]}, b^{[l]}, Z^{[l]})$

- 5) `L_model_forward(X, parameters)`:

Given training data X and parameters, `L_model_forward(X,parameters)` instantiates `linear_activation_forward()` $L-1$ times with `activation="relu"` to build the $L-1$ hidden layers, and instantiate `linear_activation_forward()` once with `activation="sigmoid"` to build the output layer. The function returns the output $A^{[L]}$ of the output layer, and caches all the intermediate results and parameters for all layers, $(A_{\text{prev}}, W, b, Z)$.

- 6) Cost function

`compute_cost(AL, Y)`: compute the cost

$$J(W, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \ln(\hat{y}^{(i)}) + (1 - y^{(i)}) \ln(1 - \hat{y}^{(i)})]$$

by

$$\text{cost} = (-1/m) * (\text{np.dot}(\text{np.log}(\text{AL}), Y.T) + \text{np.dot}(\text{np.log}(1-\text{AL}), (1-Y).T))$$

Note: the cost function is not a part of the DNN, but it outputs the cost for each iteration, and visualize the behavior of the algorithm.

- 7) `linear_backward(dZ, cache)`:

$$\begin{aligned} dW^{[l]} &= dZ^{[l]} A^{[l-1]T} * \left(\frac{1}{m}\right) & (n^{[l]}, n^{[l-1]}) &= (n^{[l]}, m) \times (n^{[l-1]}, m)^T \\ db^{[l]} &= \left(\frac{1}{m}\right) \text{np.sum}(dZ^{[l]}, \text{axis} = 1, \text{keepdims} = \text{True}) & (n^{[l]}, 1) &= (n^{[l]}, 1) \\ dA^{[l-1]} &= W^{[l]T} dZ^{[l]} & (n^{[l-1]}, 1) &= (n^{[l]}, n^{[l-1]})^T \times (n^{[l]}, 1) \end{aligned}$$

Implement the linear portion of backward propagation for a single layer (layer l)

Arguments:

`dZ` -- Gradient of the cost with respect to the linear output (of current layer l)

`cache` -- tuple of values (A_{prev}, W, b) coming from the forward propagation in the current layer

Returns:

`dA_prev` -- Gradient of the cost with respect to the activation (of the previous layer $l-1$), same shape as A_{prev}

`dW` -- Gradient of the cost with respect to W (current layer l), same shape as W

`db` -- Gradient of the cost with respect to b (current layer l), same shape as b

- 8) `linear_activation_backward(dA, cache, activation)`:

$$dZ^{[l]} = dA^{[l]} * g^{[l]'}(Z^{[l]}) \quad (n^{[l]}, m) = (n^{[l]}, m) * (n^{[l]}, m)$$

$$\begin{aligned}
dW^{[l]} &= dZ^{[l]} A^{[l-1]T} * \left(\frac{1}{m}\right) & (n^{[l]}, n^{[l-1]}) &= (n^{[l]}, m) \times (n^{[l-1]}, m)^T \\
db^{[l]} &= \left(\frac{1}{m}\right) np.sum(dZ^{[l]}, axis = 1, keepdims = True) & (n^{[l]}, 1) &= (n^{[l]}, 1) \\
dA^{[l-1]} &= W^{[l]T} dZ^{[l]} & (n^{[l-1]}, 1) &= (n^{[l]}, n^{[l-1]})^T \times (n^{[l]}, 1)
\end{aligned}$$

Implement the backward propagation for the LINEAR plus ACTIVATION layer.

Arguments:

dA -- post-activation gradient for current layer l
 cache -- tuple of values (linear_cache, activation_cache) we store for computing backward propagation efficiently
 activation -- the activation to be used in this layer, stored as a text string: "sigmoid" or "relu"

Returns:

dA_prev -- Gradient of the cost with respect to the activation (of the previous layer l-1), same shape as A_prev
 dW -- Gradient of the cost with respect to W (current layer l), same shape as W
 db -- Gradient of the cost with respect to b (current layer l), same shape as b

- 9) L_model_backward(AL, Y, caches): by instantiating linear_activation_backward(dA, cache, activation) L times, implement the backward propagation for the [LINEAR plus RELU] * (L-1) layers and LINEAR plus SIGMOID layer

Arguments:

AL -- probability vector, output of the forward propagation (L_model_forward())
 Y -- true "label" vector (containing 0 if non-cat, 1 if cat)
 caches -- list of caches containing:
 every cache of linear_activation_forward() with "relu" (it's caches[l], for l in range(L-1) i.e l = 0...L-2)
 the cache of linear_activation_forward() with "sigmoid" (it's caches[L-1])

Returns:

grads -- A dictionary with the gradients
 grads["dA" + str(l)] = ...
 grads["dW" + str(l)] = ...
 grads["db" + str(l)] = ...

- 10) update_parameters(parameters, grads, learning_rate): Update parameters using gradient descent

Arguments:

parameters -- python dictionary containing your parameters
 grads -- python dictionary containing your gradients, output of L_model_backward

Returns:

parameters -- python dictionary containing your updated parameters
 parameters["W" + str(l)] = ...
 parameters["b" + str(l)] = ...

11) predict(X, y, parameters): This function is used to predict the results of a L-layer neural network.

Arguments:

X -- data set of examples you would like to label
parameters -- parameters of the trained model

Returns:

p -- predictions for the given dataset X

12) load_data(): read the images from *.h files into arrays.

13) print_mislabeled_images(classes, X, y, p): Plots images where predictions and truth were different.

X -- dataset y -- true labels p -- predictions

dnn_utils.py

```
1.  # -*- coding: utf-8 -*-
2.  """
3.  Created on Tue Jun  4 14:26:51 2019
4.  from my assignment notes: dnn_utils.py
5.  @author: weido
6.  """
7.
8.  import numpy as np
9.  import h5py
10. import matplotlib.pyplot as plt
11. #from testCases_v4 import *
12. #from dnn_utils_v2 import sigmoid, sigmoid_backward, relu, relu_backward
13.
14. def sigmoid(Z):
15.     """
16.     Implements the sigmoid activation in numpy
17.
18.     Arguments:
19.     Z -- numpy array of any shape
20.
21.     Returns:
22.     A -- output of sigmoid(z), same shape as Z
23.     cache -- returns Z as well, useful during backpropagation
24.     """
25.
26.     A = 1/(1+np.exp(-Z))
27.     cache = Z
28.
29.     return A, cache
30.
31. def relu(Z):
32.     """
33.     Implement the RELU function.
34.     Arguments:
35.     Z -- Output of the linear layer, of any shape
36.     Returns:
37.     A -- Post-activation parameter, of the same shape as Z
38.     cache -
39.     - a python dictionary containing "A" ; stored for computing the backward pass efficiently
40.     """
```



```

40.
41.     A = np.maximum(0,Z)
42.
43.     assert(A.shape == Z.shape)
44.
45.     cache = Z
46.     return A, cache
47.
48.
49. def relu_backward(dA, Z):
50.     """
51.     Implement the backward propagation for a single RELU unit.
52.     Arguments:
53.     dA -- post-activation gradient, of any shape
54.     cache -- 'Z' where we store for computing backward propagation efficiently
55.     Returns:
56.     dZ -- Gradient of the cost with respect to Z
57.     """
58.
59.     dZ = np.array(dA, copy=True) # just converting dz to a correct object.
60.
61.     # When z <= 0, you should set dz to 0 as well.
62.     dZ[Z <= 0] = 0
63.
64.     assert (dZ.shape == Z.shape)
65.
66.     return dZ
67.
68. def sigmoid_backward(dA, Z):
69.     """
70.     Implement the backward propagation for a single SIGMOID unit.
71.     Arguments:
72.     dA -- post-activation gradient, of any shape
73.     cache -- 'Z' where we store for computing backward propagation efficiently
74.     Returns:
75.     dZ -- Gradient of the cost with respect to Z
76.     """
77.
78.     s = 1/(1+np.exp(-Z))
79.     dZ = dA * s * (1-s)
80.
81.     assert (dZ.shape == Z.shape)
82.
83.     return dZ
84.
85.
86.     %matplotlib inline
87.     plt.rcParams['figure.figsize'] = (5.0, 4.0) # set default size of plots
88.     plt.rcParams['image.interpolation'] = 'nearest'
89.     plt.rcParams['image.cmap'] = 'gray'
90.
91.     %load_ext autoreload
92.     %autoreload 2
93.
94.     np.random.seed(1)
95.
96.
97.     # FUNCTION: initialize_parameters
98.
99.     def initialize_parameters(n_x, n_h, n_y):
100.         """

```

```

101.     Argument:
102.     n_x -- size of the input layer
103.     n_h -- size of the hidden layer
104.     n_y -- size of the output layer
105.
106.     Returns:
107.     parameters -- python dictionary containing your parameters:
108.                   W1 -- weight matrix of shape (n_h, n_x)
109.                   b1 -- bias vector of shape (n_h, 1)
110.                   W2 -- weight matrix of shape (n_y, n_h)
111.                   b2 -- bias vector of shape (n_y, 1)
112.     """
113.
114.     np.random.seed(1)
115.
116.     ### START CODE HERE ### (≈ 4 lines of code)
117.     W1 = np.random.randn(n_h,n_x)/np.sqrt(n_x)*0.01
118.     b1 = np.zeros((n_h,1))
119.     W2 = np.random.randn(n_y,n_h)/np.sqrt(n_h)*0.01
120.     b2 = np.zeros((n_y,1))
121.     ### END CODE HERE ###
122.
123.     assert(W1.shape == (n_h, n_x))
124.     assert(b1.shape == (n_h, 1))
125.     assert(W2.shape == (n_y, n_h))
126.     assert(b2.shape == (n_y, 1))
127.
128.     parameters = {"W1": W1,
129.                  "b1": b1,
130.                  "W2": W2,
131.                  "b2": b2}
132.
133.     return parameters
134.
135. # FUNCTION: initialize_parameters_deep
136.
137. def initialize_parameters_deep(layer_dims):
138.     """
139.     Arguments:
140.     layer_dims -
141.     - python array (list) containing the dimensions of each layer in our network
142.
143.     Returns:
144.     parameters -
145.     - python dictionary containing your parameters "W1", "b1", ..., "WL", "bL":
146.       W1 -- weight matrix of shape (layer_dims[l], layer_dims[l-
147.       1])
148.       b1 -- bias vector of shape (layer_dims[l], 1)
149.     """
150.
151.     np.random.seed(1)
152.     parameters = {}
153.     L = len(layer_dims)            # number of layers in the network
154.
155.     for l in range(1, L):
156.         ### START CODE HERE ### (≈ 2 lines of code)
157.         parameters['W' + str(l)] = np.random.randn(layer_dims[l],layer_dims[l-

```

```

158.         assert(parameters['W' + str(l)].shape == (layer_dims[l], layer_dims[l-
159.             1]))
160.         assert(parameters['b' + str(l)].shape == (layer_dims[l], 1))
161.
162.     return parameters
163.
164. # FUNCTION: linear_forward
165.
166. def linear_forward(A, W, b):
167.     """
168.     Implement the linear part of a layer's forward propagation.
169.
170.     Arguments:
171.     A -
172.         - activations from previous layer (or input data): (size of previous layer, num
173.         ber of examples)
174.     W -
175.         - weights matrix: numpy array of shape (size of current layer, size of previous
176.         layer)
177.     b -- bias vector, numpy array of shape (size of the current layer, 1)
178.
179.     Returns:
180.     Z -- the input of the activation function, also called pre-
181.     activation parameter
182.     cache -
183.         - a python dictionary containing "A", "W" and "b" ; stored for computing the bac
184.         kward pass efficiently
185.     """
186.
187.     ### START CODE HERE ### (~ 1 line of code)
188.     Z = np.dot(W,A)+b
189.     ### END CODE HERE ###
190.
191.     assert(Z.shape == (W.shape[0], A.shape[1]))
192.     cache = (A, W, b)
193.
194.     return Z, cache
195.
196. # FUNCTION: linear_activation_forward
197.
198. def linear_activation_forward(A_prev, W, b, activation):
199.     """
200.     Implement the forward propagation for the LINEAR->ACTIVATION layer
201.
202.     Arguments:
203.     A_prev -
204.         - activations from previous layer (or input data): (size of previous layer, num
205.         ber of examples)
206.     W -
207.         - weights matrix: numpy array of shape (size of current layer, size of previous
208.         layer)
209.     b -- bias vector, numpy array of shape (size of the current layer, 1)
210.     activation -
211.         - the activation to be used in this layer, stored as a text string: "sigmoid" or
212.         "relu"
213.
214.     Returns:
215.     A -- the output of the activation function, also called the post-
216.     activation value

```

```

203.     cache -
204.     - a python dictionary containing "linear_cache" and "activation_cache";
205.       stored for computing the backward pass efficiently
206.
207.     if activation == "sigmoid":
208.         # Inputs: "A_prev, W, b". Outputs: "A, activation_cache".
209.         ### START CODE HERE ### (≈ 2 lines of code)
210.         Z, linear_cache = linear_forward(A_prev, W, b)
211.         A, activation_cache = sigmoid(Z)
212.         #A, activation_cache = A, activation_cache = sigmoid(Z)
213.         ### END CODE HERE ###
214.
215.     elif activation == "relu":
216.         # Inputs: "A_prev, W, b". Outputs: "A, activation_cache".
217.         ### START CODE HERE ### (≈ 2 lines of code)
218.         Z, linear_cache = linear_forward(A_prev, W, b)
219.         #A, activation_cache = A, activation_cache = relu(Z)
220.         A, activation_cache = relu(Z)
221.         ### END CODE HERE ###
222.
223.     assert (A.shape == (W.shape[0], A_prev.shape[1]))
224.     cache = (linear_cache, activation_cache)
225.
226.     return A, cache
227.
228. # FUNCTION: L_model_forward
229.
230. def L_model_forward(X, parameters):
231.     """
232.     Implement forward propagation for the [LINEAR->RELU]*(L-1)->LINEAR-
233.     >SIGMOID computation
234.
235.     Arguments:
236.     X -- data, numpy array of shape (input size, number of examples)
237.     parameters -- output of initialize_parameters_deep()
238.
239.     Returns:
240.     AL -- last post-activation value
241.     caches -- list of caches containing:
242.         every cache of linear_activation_forward() (there are L-
243.         1 of them, indexed from 0 to L-1)
244.     """
245.
246.     caches = []
247.     A = X
248.     L = len(parameters) // 2 # number of layers in the neural network
249.
250.     # Implement [LINEAR -> RELU]*(L-1). Add "cache" to the "caches" list.
251.     for l in range(1, L):
252.         A_prev = A
253.         ### START CODE HERE ### (≈ 2 lines of code)
254.         A, cache = linear_activation_forward(A_prev, parameters['W'+str(l)], parameters['b'+str(l)], activation = "relu")
255.         caches.append(cache)
256.         ### END CODE HERE ###
257.
258.     # Implement LINEAR -> SIGMOID. Add "cache" to the "caches" list.
259.     ### START CODE HERE ### (≈ 2 lines of code)

```

```

258.     # A, linear_activation_cache = linear_activation_forward(A_prev, W, b, activation = "sigmoid")
259.     AL, cache = linear_activation_forward(A, parameters['W'+str(L)], parameters['b'+str(L)], activation = "sigmoid")
260.     caches.append(cache)
261.     ### END CODE HERE ###
262.
263.     assert(AL.shape == (1,X.shape[1]))
264.
265.     return AL, caches
266.
267.
268. # FUNCTION: compute_cost
269.
270. def compute_cost(AL, Y):
271.     """
272.     Implement the cost function defined by equation (7).
273.
274.     Arguments:
275.     AL -
276.     - probability vector corresponding to your label predictions, shape (1, number of examples)
277.     Y -- true "label" vector (for example: containing 0 if non-cat, 1 if cat), shape (1, number of examples)
278.
279.     Returns:
280.     cost -- cross-entropy cost
281.
282.     """
283.
284.     m = Y.shape[1]
285.
286.     # Compute loss from aL and y.
287.     ### START CODE HERE ### (~ 1 lines of code)
288.     cost = (-1/m)*(np.dot(np.log(AL), Y.T)+np.dot(np.log(1-AL), (1-Y).T))
289.     ### END CODE HERE ###
290.
291.     cost = np.squeeze(cost)      # To make sure your cost's shape is what we expect (e.g. this turns [[17]] into 17).
292.     assert(cost.shape == ())
293.
294.     return cost
295.
296. # FUNCTION: linear_backward
297.
298. def linear_backward(dZ, cache):
299.     """
300.     Implement the linear portion of backward propagation for a single layer (layer l)
301.
302.     Arguments:
303.     dZ -
304.     - Gradient of the cost with respect to the linear output (of current layer l)
305.     cache -
306.     - tuple of values (A_prev, W, b) coming from the forward propagation in the current layer
307.
308.     Returns:
309.     dA_prev -
310.     - Gradient of the cost with respect to the activation (of the previous layer l-1), same shape as A_prev

```

```

306.     dW -
        - Gradient of the cost with respect to W (current layer l), same shape as W
307.     db -
        - Gradient of the cost with respect to b (current layer l), same shape as b
308.     """
309.     A_prev, W, b = cache
310.     m = A_prev.shape[1]
311.
312.     ### START CODE HERE ### (~ 3 lines of code)
313.     dW = (1/m)*np.dot(dZ, A_prev.T)
314.     db = (1/m)*np.sum(dZ, axis=1, keepdims = True)
315.     dA_prev = np.dot(W.T, dZ)
316.     ### END CODE HERE ###
317.
318.     assert (dA_prev.shape == A_prev.shape)
319.     assert (dW.shape == W.shape)
320.     assert (db.shape == b.shape)
321.
322.     return dA_prev, dW, db
323.
324. # FUNCTION: linear_activation_backward
325.
326. def linear_activation_backward(dA, cache, activation):
327.     """
328.     Implement the backward propagation for the LINEAR->ACTIVATION layer.
329.
330.     Arguments:
331.     dA -- post-activation gradient for current layer l
332.     cache -
        - tuple of values (linear_cache, activation_cache) we store for computing backward propagation efficiently
333.     activation -
        - the activation to be used in this layer, stored as a text string: "sigmoid" or "relu"
334.
335.     Returns:
336.     dA_prev -
        - Gradient of the cost with respect to the activation (of the previous layer l-1), same shape as A_prev
337.     dW -
        - Gradient of the cost with respect to W (current layer l), same shape as W
338.     db -
        - Gradient of the cost with respect to b (current layer l), same shape as b
339.     """
340.     linear_cache, activation_cache = cache
341.
342.     if activation == "relu":
343.         ### START CODE HERE ### (~ 2 lines of code)
344.         dZ = relu_backward(dA, activation_cache)
345.         dA_prev, dW, db = linear_backward(dZ, linear_cache)
346.         ### END CODE HERE ###
347.
348.     elif activation == "sigmoid":
349.         ### START CODE HERE ### (~ 2 lines of code)
350.         dZ = sigmoid_backward(dA, activation_cache)
351.         dA_prev, dW, db = linear_backward(dZ, linear_cache)
352.         ### END CODE HERE ###
353.
354.     return dA_prev, dW, db
355.
356. # FUNCTION: L_model_backward

```

```

357.
358. def L_model_backward(AL, Y, caches):
359.     """
360.     Implement the backward propagation for the [LINEAR->RELU] * (L-1) -
    > LINEAR -> SIGMOID group
361.
362.     Arguments:
363.     AL -
    - probability vector, output of the forward propagation (L_model_forward())
364.     Y -- true "label" vector (containing 0 if non-cat, 1 if cat)
365.     caches -- list of caches containing:
366.         every cache of linear_activation_forward() with "relu" (it's caches[1],
    for l in range(L-1) i.e l = 0...L-2)
367.         the cache of linear_activation_forward() with "sigmoid" (it's caches[L-1])
368.
369.     Returns:
370.     grads -- A dictionary with the gradients
371.         grads["dA" + str(l)] = ...
372.         grads["dW" + str(l)] = ...
373.         grads["db" + str(l)] = ...
374.     """
375.     grads = {}
376.     L = len(caches) # the number of layers
377.     m = AL.shape[1]
378.     Y = Y.reshape(AL.shape) # after this line, Y is the same shape as AL
379.
380.     # Initializing the backpropagation
381.     ### START CODE HERE ### (1 line of code)
382.     dAL = -(np.divide(Y,AL)-np.divide(1-Y,1-AL))
383.     ### END CODE HERE ###
384.
385.     # Lth layer (SIGMOID -
    > LINEAR) gradients. Inputs: "dAL, current_cache". Outputs: "grads["dA-
    1"], grads["dWL"], grads["dbL"]
386.     ### START CODE HERE ### (approx. 2 lines)
387.     current_cache = caches[L-1]
388.     grads["dA" + str(L-
    1)], grads["dW" + str(L)], grads["db" + str(L)] = linear_activation_backward(dAL
    , current_cache, "sigmoid")
389.     ### END CODE HERE ###
390.
391.     # Loop from l=L-2 to l=0
392.     for l in reversed(range(L-1)):
393.         # lth layer: (RELU -> LINEAR) gradients.
394.         # Inputs: "grads["dA" + str(l + 1)], current_cache". Outputs: "grads["dA
    " + str(l)] , grads["dW" + str(l + 1)] , grads["db" + str(l + 1)]
395.         ### START CODE HERE ### (approx. 5 lines)
396.         current_cache = caches[l]
397.         dA_prev_temp, dW_temp, db_temp = linear_activation_backward(grads["dA" + s
    tr(l+1)], current_cache, "relu")
398.         grads["dA" + str(l)] = dA_prev_temp
399.         grads["dW" + str(l + 1)] = dW_temp
400.         grads["db" + str(l + 1)] = db_temp
401.         ### END CODE HERE ###
402.
403.     return grads
404.
405. # FUNCTION: update_parameters
406.
407. def update_parameters(parameters, grads, learning_rate):

```

```

408.     """
409.     Update parameters using gradient descent
410.
411.     Arguments:
412.     parameters -- python dictionary containing your parameters
413.     grads -
414.         - python dictionary containing your gradients, output of L_model_backward
415.
416.     Returns:
417.     parameters -- python dictionary containing your updated parameters
418.         parameters["W" + str(l)] = ...
419.         parameters["b" + str(l)] = ...
420.     """
421.     L = len(parameters) // 2 # number of layers in the neural network
422.
423.     # Update rule for each parameter. Use a for loop.
424.     ### START CODE HERE ### (~ 3 lines of code)
425.     for l in range(L):
426.         parameters["W" + str(l+1)] = parameters["W" + str(l+1)] -
427.             learning_rate*grads["dw" + str(l+1)]
428.         parameters["b" + str(l+1)] = parameters["b" + str(l+1)] -
429.             learning_rate*grads["db" + str(l+1)]
430.     ### END CODE HERE ###
431.     return parameters
432.
433. def load_data():
434.     train_dataset = h5py.File('train_catvnoncat.h5', "r")
435.     train_set_x_orig = np.array(train_dataset["train_set_x"][:]) # your train set features
436.     train_set_y_orig = np.array(train_dataset["train_set_y"][:]) # your train set labels
437.
438.     test_dataset = h5py.File('test_catvnoncat.h5', "r")
439.     test_set_x_orig = np.array(test_dataset["test_set_x"][:]) # your test set features
440.     test_set_y_orig = np.array(test_dataset["test_set_y"][:]) # your test set labels
441.
442.     classes = np.array(test_dataset["list_classes"][:]) # the list of classes
443.
444.     train_set_y_orig = train_set_y_orig.reshape((1, train_set_y_orig.shape[0]))
445.     test_set_y_orig = test_set_y_orig.reshape((1, test_set_y_orig.shape[0]))
446.
447.     return train_set_x_orig, train_set_y_orig, test_set_x_orig, test_set_y_orig, classes
448.
449. def predict(X, y, parameters):
450.     """
451.     This function is used to predict the results of a L-layer neural network.
452.
453.     Arguments:
454.     X -- data set of examples you would like to label
455.     parameters -- parameters of the trained model
456.
457.     Returns:
458.     p -- predictions for the given dataset X
459.     """
460.     m = X.shape[1]

```



```

460.     n = len(parameters) // 2 # number of layers in the neural network
461.     p = np.zeros((1,m))
462.
463.     # Forward propagation
464.     probas, caches = L_model_forward(X, parameters)
465.
466.
467.     # convert probas to 0/1 predictions
468.     for i in range(0, probas.shape[1]):
469.         if probas[0,i] > 0.5:
470.             p[0,i] = 1
471.         else:
472.             p[0,i] = 0
473.
474.     #print results
475.     #print ("predictions: " + str(p))
476.     #print ("true labels: " + str(y))
477.     print("Accuracy: " + str(np.sum((p == y)/m)))
478.
479.     return p
480.
481. def print_mislabeled_images(classes, X, y, p):
482.     """
483.     Plots images where predictions and truth were different.
484.     X -- dataset
485.     y -- true labels
486.     p -- predictions
487.     """
488.     a = p + y
489.     mislabeled_indices = np.asarray(np.where(a == 1))
490.     plt.rcParams['figure.figsize'] = (40.0, 40.0) # set default size of plots
491.     num_images = len(mislabeled_indices[0])
492.     for i in range(num_images):
493.         index = mislabeled_indices[1][i]
494.
495.         plt.subplot(2, num_images, i + 1)
496.         plt.imshow(X[:,index].reshape(64,64,3), interpolation='nearest')
497.         plt.axis('off')
498.         plt.title("Prediction: " + classes[int(p[0,index])].decode("utf-
8") + " \n Class: " + classes[y[0,index]].decode("utf-8"))

```

7.3.2 Package (h5py and pillow) and dataset preparation

Before we develop the python code for the DNN, we need to import the following packages:

- h5py is a common package to interact with a dataset that is stored on an H5 file.
- PIL and scipy are used here to test the DNN model with our own picture at the end.

1) Install h5py

```
(base) C:\Users\weido>activate tensorflow_cpu
```

```
(tensorflow_cpu) C:\Users\weido>conda install h5py
```

```
Collecting package metadata: done
```

```
Solving environment: done
```

```
## Package Plan ##
```

```
environment location: C:\Users\weido\AppData\Local\Continuum\anaconda3\envs\tensorflow_cpu
```

```
added / updated specs:
```

```
- h5py
```

The following packages will be downloaded:

package	build	
h5py-2.9.0	py36h5e291fa_0	969 KB
hdf5-1.10.4	h7ebc959_0	19.2 MB
intel-openmp-2019.4	245	1.7 MB
mkl-2019.4	245	157.5 MB
mkl_fft-1.0.12	py36h14836fe_0	136 KB
mkl_random-1.0.2	py36h343c172_0	318 KB
numpy-1.16.4	py36h19fb1c0_0	49 KB
numpy-base-1.16.4	py36hc3f5095_0	4.1 MB
openssl-1.1.1c	he774522_1	5.7 MB
pyreadline-2.1	py36_1	141 KB
Total:		189.9 MB

The following NEW packages will be INSTALLED:

blas	pkgs/main/win-64::blas-1.0-mkl
h5py	pkgs/main/win-64::h5py-2.9.0-py36h5e291fa_0
hdf5	pkgs/main/win-64::hdf5-1.10.4-h7ebc959_0
icc_rt	pkgs/main/win-64::icc_rt-2019.0.0-h0cc432a_1
intel-openmp	pkgs/main/win-64::intel-openmp-2019.4-245
mkl	pkgs/main/win-64::mkl-2019.4-245
mkl_fft	pkgs/main/win-64::mkl_fft-1.0.12-py36h14836fe_0
mkl_random	pkgs/main/win-64::mkl_random-1.0.2-py36h343c172_0
numpy	pkgs/main/win-64::numpy-1.16.4-py36h19fb1c0_0
numpy-base	pkgs/main/win-64::numpy-base-1.16.4-py36hc3f5095_0
pyreadline	pkgs/main/win-64::pyreadline-2.1-py36_1

The following packages will be UPDATED:

openssl	1.1.1b-he774522_1 --> 1.1.1c-he774522_1
---------	---

Proceed ([y]/n)? y

Downloading and Extracting Packages

h5py-2.9.0	969 KB	##### 100%
mkl_fft-1.0.12	136 KB	##### 100%
numpy-1.16.4	49 KB	##### 100%
mkl-2019.4	157.5 MB	##### 100%
intel-openmp-2019.4	1.7 MB	##### 100%
numpy-base-1.16.4	4.1 MB	##### 100%
openssl-1.1.1c	5.7 MB	##### 100%
mkl_random-1.0.2	318 KB	##### 100%
pyreadline-2.1	141 KB	##### 100%
hdf5-1.10.4	19.2 MB	##### 100%

Preparing transaction: done
Verifying transaction: done
Executing transaction: done

2) Install PIL

(tensorflow_cpu) C:\Users\weido>conda install pillow

Collecting package metadata: done

Solving environment: done

Package Plan

environment location: C:\Users\weido\AppData\Local\Continuum\anaconda3\envs\tensorflow_cpu

added / updated specs:
- pillow

The following packages will be downloaded:

package	build	
ca-certificates-2019.5.15	0	166 KB
libtiff-4.0.10	hb898794_2	1.1 MB
olefile-0.46	py36_0	49 KB
pillow-6.0.0	py36hdc69c19_0	699 KB
Total:		1.9 MB

The following NEW packages will be INSTALLED:

```
freetype      pkgs/main/win-64::freetype-2.9.1-ha9979f8_1
libtiff       pkgs/main/win-64::libtiff-4.0.10-hb898794_2
olefile       pkgs/main/win-64::olefile-0.46-py36_0
pillow        pkgs/main/win-64::pillow-6.0.0-py36hdc69c19_0
tk            pkgs/main/win-64::tk-8.6.8-hfa6e2cd_0
xz            pkgs/main/win-64::xz-5.2.4-h2fa13f4_4
zstd          pkgs/main/win-64::zstd-1.3.7-h508b16e_0
```

The following packages will be UPDATED:

```
ca-certificates      2019.1.23-0 --> 2019.5.15-0
```

Proceed ([y]/n)? y

Downloading and Extracting Packages

```
libtiff-4.0.10 | 1.1 MB | ##### | 100%
pillow-6.0.0   | 699 KB | ##### | 100%
olefile-0.46   | 49 KB | ##### | 100%
ca-certificates-2019 | 166 KB | ##### | 100%
```

Preparing transaction: done

Verifying transaction: done

Executing transaction: done

- 3) Download train_catvnoncat.h5 and test_catvnoncat.h5. These two files will be used for DNN training and testing.

7.3.3 Put all together (week4_assign.py)

Top-level file “week4_assign.py” is pasted below and followed by explanations.

```
1.  # -*- coding: utf-8 -*-
2.  """
3.  Created on Tue Jun  4 15:51:29 2019
4.  week4_assign.py
5.  @author: weido
6.  """
7.
8.  import time
9.  import numpy as np
10. import h5py
11. import matplotlib.pyplot as plt
12. import scipy
13. from PIL import Image
14. from scipy import ndimage
15. #from dnn_app_utils_v3 import *
```

```

16. from dnn_utils import *
17.
18. %matplotlib inline
19. plt.rcParams['figure.figsize'] = (5.0, 4.0) # set default size of plots
20. plt.rcParams['image.interpolation'] = 'nearest'
21. plt.rcParams['image.cmap'] = 'gray'
22.
23. %load_ext autoreload
24. %autoreload 2
25.
26. np.random.seed(1)
27.
28. train_x_orig, train_y, test_x_orig, test_y, classes = load_data()
29. # Example of a picture
30. index = 10
31. plt.imshow(train_x_orig[index])
32. plt.show()
33. print ("y = " + str(train_y[0,index]) + ". It's a " + classes[train_y[0,index]].
    decode("utf-8") + " picture.")
34.
35. # Example of a picture
36. index = 11
37. plt.imshow(train_x_orig[index])
38. plt.show()
39. print ("y = " + str(train_y[0,index]) + ". It's a " + classes[train_y[0,index]].
    decode("utf-8") + " picture.")
40.
41. # Explore your dataset
42. m_train = train_x_orig.shape[0]
43. num_px = train_x_orig.shape[1]
44. m_test = test_x_orig.shape[0]
45.
46. print ("Number of training examples: " + str(m_train))
47. print ("Number of testing examples: " + str(m_test))
48. print ("Each image is of size: (" + str(num_px) + ", " + str(num_px) + ", 3)")
49. print ("train_x_orig shape: " + str(train_x_orig.shape))
50. print ("train_y shape: " + str(train_y.shape))
51. print ("test_x_orig shape: " + str(test_x_orig.shape))
52. print ("test_y shape: " + str(test_y.shape))
53. # Reshape the training and test examples
54. train_x_flatten = train_x_orig.reshape(train_x_orig.shape[0], -1).T # The "-
    1" makes reshape flatten the remaining dimensions
55. test_x_flatten = test_x_orig.reshape(test_x_orig.shape[0], -1).T
56.
57. # Standardize data to have feature values between 0 and 1.
58. train_x = train_x_flatten/255.
59. test_x = test_x_flatten/255.
60.
61. print ("train_x's shape: " + str(train_x.shape))
62. print ("test_x's shape: " + str(test_x.shape))
63.
64. #### CONSTANTS DEFINING THE MODEL ####
65. n_x = 12288 # num_px * num_px * 3
66. n_h = 7
67. n_y = 1
68. layers_dims = (n_x, n_h, n_y)
69.
70.
71.
72. # GRADED FUNCTION: two_layer_model
73.

```

```

74. def two_layer_model(X, Y, layers_dims, learning_rate = 0.0075, num_iterations =
    3000, print_cost=False):
75.     """
76.     Implements a two-layer neural network: LINEAR->RELU->LINEAR->SIGMOID.
77.
78.     Arguments:
79.     X -- input data, of shape (n_x, number of examples)
80.     Y -- true "label" vector (containing 0 if cat, 1 if non-
        cat), of shape (1, number of examples)
81.     layers_dims -- dimensions of the layers (n_x, n_h, n_y)
82.     num_iterations -- number of iterations of the optimization loop
83.     learning_rate -- learning rate of the gradient descent update rule
84.     print_cost -
        - If set to True, this will print the cost every 100 iterations
85.
86.     Returns:
87.     parameters -- a dictionary containing W1, W2, b1, and b2
88.     """
89.
90.     np.random.seed(1)
91.     grads = {}
92.     costs = [] # to keep track of the cost
93.     m = X.shape[1] # number of examples
94.     (n_x, n_h, n_y) = layers_dims
95.
96.     # Initialize parameters dictionary, by calling one of the functions you'd pr
        eviously implemented
97.     ### START CODE HERE ### (~ 1 line of code)
98.     parameters = initialize_parameters(n_x, n_h, n_y)
99.     ### END CODE HERE ###
100.
101.     # Get W1, b1, W2 and b2 from the dictionary parameters.
102.     W1 = parameters["W1"]
103.     b1 = parameters["b1"]
104.     W2 = parameters["W2"]
105.     b2 = parameters["b2"]
106.
107.     # Loop (gradient descent)
108.
109.     for i in range(0, num_iterations):
110.
111.         # Forward propagation: LINEAR -> RELU -> LINEAR -
            > SIGMOID. Inputs: "X, W1, b1, W2, b2". Output: "A1, cache1, A2, cache2".
112.         ### START CODE HERE ### (~ 2 lines of code)
113.         A1, cache1 = linear_activation_forward(X, W1, b1, "relu")
114.         A2, cache2 = linear_activation_forward(A1, W2, b2, "sigmoid")
115.         ### END CODE HERE ###
116.
117.         # Compute cost
118.         ### START CODE HERE ### (~ 1 line of code)
119.         cost = compute_cost(A2, Y)
120.         ### END CODE HERE ###
121.
122.         # Initializing backward propagation
123.         dA2 = - (np.divide(Y, A2) - np.divide(1 - Y, 1 - A2))
124.
125.         # Backward propagation. Inputs: "dA2, cache2, cache1". Outputs: "dA1, dW
            2, db2; also dA0 (not used), dW1, db1".
126.         ### START CODE HERE ### (~ 2 lines of code)
127.         dA1, dW2, db2 = linear_activation_backward(dA2, cache2, "sigmoid")
128.         dA0, dW1, db1 = linear_activation_backward(dA1, cache1, "relu")

```

```

129.         ### END CODE HERE ###
130.
131.         # Set grads['dW1'] to dW1, grads['db1'] to db1, grads['dW2'] to dW2, gra
ds['db2'] to db2
132.         grads['dW1'] = dW1
133.         grads['db1'] = db1
134.         grads['dW2'] = dW2
135.         grads['db2'] = db2
136.
137.         # Update parameters.
138.         ### START CODE HERE ### (approx. 1 line of code)
139.         parameters = update_parameters(parameters, grads, learning_rate)
140.         ### END CODE HERE ###
141.
142.         # Retrieve W1, b1, W2, b2 from parameters
143.         W1 = parameters["W1"]
144.         b1 = parameters["b1"]
145.         W2 = parameters["W2"]
146.         b2 = parameters["b2"]
147.
148.         # Print the cost every 100 training example
149.         if print_cost and i % 100 == 0:
150.             print("Cost after iteration {}: {}".format(i, np.squeeze(cost)))
151.         if print_cost and i % 100 == 0:
152.             costs.append(cost)
153.
154.         # plot the cost
155.
156.         plt.plot(np.squeeze(costs))
157.         plt.ylabel('cost')
158.         plt.xlabel('iterations (per tens)')
159.         plt.title("Learning rate = " + str(learning_rate))
160.         plt.show()
161.
162.         return parameters
163.
164.
165.     parameters = two_layer_model(train_x, train_y, layers_dims = (n_x, n_h, n_y), nu
m_iterations = 2500, print_cost=True)
166.
167.     predictions_train = predict(train_x, train_y, parameters)
168.
169.     predictions_test = predict(test_x, test_y, parameters)
170.
171.     ### CONSTANTS ###
172.     layers_dims = [12288, 20, 7, 5, 1] # 4-layer model
173.
174.     # GRADED FUNCTION: L_layer_model
175.
176.     def L_layer_model(X, Y, layers_dims, learning_rate = 0.0075, num_iterations = 30
00, print_cost=False):#lr was 0.009
177.         """
178.         Implements a L-layer neural network: [LINEAR->RELU]*(L-1)->LINEAR-
>SIGMOID.
179.
180.         Arguments:
181.         X -- data, numpy array of shape (number of examples, num_px * num_px * 3)
182.         Y -- true "label" vector (containing 0 if cat, 1 if non-
cat), of shape (1, number of examples)

```

```

183.     layers_dims -
        - list containing the input size and each layer size, of length (number of layers + 1).
184.     learning_rate -- learning rate of the gradient descent update rule
185.     num_iterations -- number of iterations of the optimization loop
186.     print_cost -- if True, it prints the cost every 100 steps
187.
188.     Returns:
189.     parameters -
        - parameters learnt by the model. They can then be used to predict.
190.     """
191.
192.     np.random.seed(1)
193.     costs = [] # keep track of cost
194.
195.     # Parameters initialization. (~ 1 line of code)
196.     """ START CODE HERE """
197.     parameters = initialize_parameters_deep(layers_dims)
198.     """ END CODE HERE """
199.
200.     # Loop (gradient descent)
201.     for i in range(0, num_iterations):
202.
203.         # Forward propagation: [LINEAR -> RELU]*(L-1) -> LINEAR -> SIGMOID.
204.         """ START CODE HERE """ (~ 1 line of code)
205.         AL, caches = L_model_forward(X, parameters)
206.         """ END CODE HERE """
207.
208.         # Compute cost.
209.         """ START CODE HERE """ (~ 1 line of code)
210.         cost = compute_cost(AL, Y)
211.         """ END CODE HERE """
212.
213.         # Backward propagation.
214.         """ START CODE HERE """ (~ 1 line of code)
215.         grads = L_model_backward(AL, Y, caches)
216.         """ END CODE HERE """
217.
218.         # Update parameters.
219.         """ START CODE HERE """ (~ 1 line of code)
220.         parameters = update_parameters(parameters, grads, learning_rate)
221.         """ END CODE HERE """
222.
223.         # Print the cost every 100 training example
224.         if print_cost and i % 100 == 0:
225.             print ("Cost after iteration %i: %f" %(i, cost))
226.         if print_cost and i % 100 == 0:
227.             costs.append(cost)
228.
229.     # plot the cost
230.     plt.plot(np.squeeze(costs))
231.     plt.ylabel('cost')
232.     plt.xlabel('iterations (per tens)')
233.     plt.title("Learning rate =" + str(learning_rate))
234.     plt.show()
235.
236.     return parameters
237.
238. parameters = L_layer_model(train_x, train_y, layers_dims, num_iterations = 2500,
        print_cost = True)
239.

```


```

240. pred_train = predict(train_x, train_y, parameters)
241.
242. #Accuracy: 0.985645933014
243.
244. pred_test = predict(test_x, test_y, parameters)
245.
246. #Accuracy: 0.8
247. print_mislabeled_images(classes, test_x, test_y, pred_test)
248.
249. import PIL
250.
251.
252. im=Image.open(r"cat13.jpg")
253. im = im.resize((64,64), PIL.Image.ANTIALIAS)
254. temp_file='sized.jpg'
255. im.save(temp_file)
256. pic = Image.open(temp_file)
257. pix = np.array(pic.getdata()).reshape(pic.size[0]*pic.size[1]*3, 1)
258. pix=pix/255.
259.
260. my_label=[1]
261. my_predicted_image = predict(pix, my_label, parameters)
262.
263. plt.imshow(im)
264. print ("y = " + str(np.squeeze(my_predicted_image)) + ", your L-
layer model predicts a \"" + classes[int(np.squeeze(my_predicted_image)),].decode("utf-8") + "\" picture.")


```

Explanations:

- 1) Lines 7-16 import all packages.
- 2) Lines 26-40 load the dataset, and display some example pictures.



y = 0. It's a non-cat picture.



y = 1. It's a cat picture.

- 3) Lines 41-63 print the sizes and shapes of the data, and prepare inputs for neural network.


```

Number of training examples: 209
Number of testing examples: 50
Each image is of size: (64, 64, 3)
train_x_orig shape: (209, 64, 64, 3)
train_y shape: (1, 209)
test_x_orig shape: (50, 64, 64, 3)
test_y shape: (1, 50)
train_x's shape: (12288, 209)
test_x's shape: (12288, 50)


```
- 4) Lines 71-162 define a 2-layer neural network.
- 5) Lines 164-169 learn the 2-layer nn model, and predict based on training set and test set.


```

Cost after iteration 0: 0.6950464961800915
Cost after iteration 100: 0.5892596054583805
Cost after iteration 200: 0.5232609173622991
Cost after iteration 300: 0.4497686396221906
Cost after iteration 400: 0.42090021618838996
Cost after iteration 500: 0.37246403061745953

```



Cost after iteration 600: 0.3474205187020189
 Cost after iteration 700: 0.31719191987370277
 Cost after iteration 800: 0.266437743477466
 Cost after iteration 900: 0.21991432807842576
 Cost after iteration 1000: 0.1435789889362377
 Cost after iteration 1100: 0.4530921262322144
 Cost after iteration 1200: 0.09499357670093511
 Cost after iteration 1300: 0.08014128076781371
 Cost after iteration 1400: 0.0694023400553646
 Cost after iteration 1500: 0.06021664023174591
 Cost after iteration 1600: 0.05327415758001876
 Cost after iteration 1700: 0.047629032620984335
 Cost after iteration 1800: 0.04297588879436867
 Cost after iteration 1900: 0.039036074365138215
 Cost after iteration 2000: 0.03568313638049026
 Cost after iteration 2100: 0.03291526373054676
 Cost after iteration 2200: 0.03047219305912061
 Cost after iteration 2300: 0.028387859212946117
 Cost after iteration 2400: 0.02661521237277608



Accuracy: 0.9999999999999998
 Accuracy: 0.74

- 6) Lines 174-236 define a L-layer nn model.
- 7) Lines 237-248 learn a particular L-layer nn model with L=4, and make predictions.
 layers_dims = [12288, 20, 7, 5, 1] # 4-layer model

Cost after iteration 0: 0.771749
 Cost after iteration 100: 0.672053
 Cost after iteration 200: 0.648263
 Cost after iteration 300: 0.611507
 Cost after iteration 400: 0.567047
 Cost after iteration 500: 0.540138
 Cost after iteration 600: 0.527930
 Cost after iteration 700: 0.465477
 Cost after iteration 800: 0.369126
 Cost after iteration 900: 0.391747
 Cost after iteration 1000: 0.315187
 Cost after iteration 1100: 0.272700
 Cost after iteration 1200: 0.237419
 Cost after iteration 1300: 0.199601
 Cost after iteration 1400: 0.189263
 Cost after iteration 1500: 0.161189
 Cost after iteration 1600: 0.148214
 Cost after iteration 1700: 0.137775
 Cost after iteration 1800: 0.129740
 Cost after iteration 1900: 0.121225
 Cost after iteration 2000: 0.113821
 Cost after iteration 2100: 0.107839
 Cost after iteration 2200: 0.102855
 Cost after iteration 2300: 0.100897
 Cost after iteration 2400: 0.092878



Accuracy: 0.9856459330143539

Accuracy: 0.8

Miss-classified images (10 images) in test dataset:



8) Lines 248-264 test on an extra jpg image.

7.3.4 Results

In this section, we developed multiple-layer neural networks for image (cat or non-cat) classification. The training set includes 209 images and testing set has 50 images. Each image is (64, 64, 3). A two-layer network with 7 units in the hidden layer achieves an accuracy of 0.74 on testing set, and a four-layer network with 20, 7, and 5 units in the three hidden layers gets an accuracy of 0.8.

Files for the project:

- dnn_utils.py: functions
- week4_assign.py: main program
- train_catvnoncat.h5: training set
- test_catvnoncat.h5: testing set

Summary

This chapter introduces ReLu activation function for image feature extraction, which is an important activation function in computer vision. We developed multiple-layer neural networks (any number of layers and any number units in each layer) from the scratch, using basic gradient descent algorithm.

References

- [1] “Neural networks and deep learning” online course at www.coursera.org

Exercises

1. Develop a 5-layer neural network to classify the images in the dataset train_catvnoncat.h5 and test_catvnoncat.h5. The number of units in the four hidden layers are 32, 16, 8, 4, respectively. All hidden layers use ReLU activation function and output layer uses sigmoid function for binary classification (cat or non-cat).
 - 1) Plot the cost function versus iterations.
 - 2) Print the prediction accuracies on training set and testing set.
2. Try different sizes by assigning different values to *layers_dims* in week4_assign.py. Can you find a value for *layers_dims* such that the neural network outperforms the neural network with *layers_dims*=[12288,20,7,5,1]?

Note: you can change both the length of *layers_dims* and the values in it.